

# Spatio-Temporal Deep Learning for Traffic Forecasting

A Comparative Study of LSTM  
and DCRNN on PEMS-BAY



# Motivation

- **Traffic congestion continues to grow worldwide**, leading to significant losses in time, fuel consumption, and economic productivity.
- **Real-time traffic prediction is critical** for enabling dynamic routing and navigation, adaptive traffic signal control, incident detection and emergency response, improved public transportation scheduling.
- Urban road networks exhibit strong **spatial and temporal** dependencies.



# Objectives & Research Questions


**Main Goal:** Evaluate how temporal-only (LSTM) vs. spatio-temporal (DCRNN) deep learning models perform on short-term traffic forecasting.

## Questions:

1. Does modeling spatial graph connectivity significantly improve forecasting accuracy?
2. How does prediction accuracy change across different forecasting horizons (5-30 minutes)?
3. What advantages does diffusion-based graph modeling offer over sequence-only models?
4. How well do LSTM and DCRNN generalize to real-world traffic patterns in PEMS-BAY?

# PEMS-BAY Overview

<b>Dataset</b>	PEMS-BAY
<b>Sensors</b>	325
<b>Period</b>	1 Jan - 30 Jun 2017
<b>Region</b>	San Francisco Bay Area, California
<b>Sampling Rate</b>	Every 5 minutes
<b>Data Type</b>	Traffic speed measurements (multivariate spatio-temporal series)



*PEMS Bay Dataset: Every red dot here in the map depicts the sensor spatially located across the freeways.*

# Data Preprocessing

## Normalization

- Applied **z-score normalization** to speed values of all sensors to scale them.
- Loaded and used the **precomputed 325×325 adjacency matrix** (distance-based graph) to encode spatial relationships between sensors.

## Sequence Construction


- Converted continuous data into **sliding windows of 12 past time steps : 1st, 2nd and 3rd future predictions**

## Dataset Split

- 70% training, 10% validation, 20% testing

## Graph Structure Visualization

- Adjacency matrix visualized using a heatmap to show connectivity strength between sensors and the spatial structure learned by DCRNN.




# Baseline Model: LSTM

**LSTM learns temporal patterns but not spatial structure.**

Each sensor is modeled independently: past speeds > future speeds.


**Flow:** Input sequence > LSTM Layers > Fully Connected Layer > Predicted Future Speeds



# Advanced Model: DCRNN

**DCRNN learns both spatial (graph connectivity) and temporal patterns.**  
Uses diffusion graph convolutions + recurrent modeling to capture how traffic moves through the network.

**Flow:** Graph > DiffusionConv > Recurrent Layers > Output.



# Training & Evaluation Setup

## Training Configuration


- **Loss Function:** Mean Squared Error (MSE)
- **Optimizer:** Adam (learning rate = 0.001)
- **Batch Size:** 64
- **Epochs:** 20
- **Hardware:** Google Colab GPU environment (Tesla T4)

## Evaluation Metrics


- **MAE** - Mean Absolute Error
- **RMSE** - Root Mean Squared Error

## Model Validation

- Training and validation losses monitored to track convergence and detect overfitting.
- LSTM and DCRNN were evaluated under identical preprocessing, sequence lengths, and dataset splits for fair comparison.



Train vs Epoch curve: LSTM





Train vs Epoch curve: DCRNN

# Experimental Results: PEMS-BAY

Training and validation curves converge for all models:

- LSTM achieves the lowest plateaued loss at each horizon
- DCRNN without adjacency converges to moderately higher loss
- Graph-based DCRNN with adjacency has the highest loss in our configuration



# Performance Metrics: PEMS-BAY

Metrics computed on held-out test set  
MSE, MAE, RMSE:

- Across 5-30 minute horizons, LSTM consistently performs best
- DCRNN without adjacency is slightly worse but still competitive
- DCRNN with adjacency has substantially larger errors at every horizon

Horizon	Model	Test MSE	Test MAE	Test RMSE
5 min	LSTM	0.0426	0.8685	1.6603
	DCRNN no-adj	0.0737	0.9841	1.8501
	DCRNN with-adj	0.2209	1.7321	3.2884
10 min	LSTM	0.0668	1.0357	2.0803
	DCRNN no-adj	0.1031	1.1601	2.2525
	DCRNN with-adj	0.2446	1.8163	3.4893
15 min	LSTM	0.0925	1.204	2.4471
	DCRNN no-adj	0.1322	1.332	2.5764
	DCRNN with-adj	0.2713	1.9151	3.6935
30 min	LSTM	0.1628	1.5399	3.2467
	DCRNN no-adj	0.2217	1.744	3.3933
	DCRNN with-adj	0.3367	2.1368	4.1964

# Contradictory findings

**Expectation:** Research suggests Spatio-Temporal Graph models like DCRNN consistently outperform temporal-only baselines like LSTM.

## Our Empirical Reality:

- **LSTM Superiority:** A well-tuned LSTM consistently achieved lower error rates across 5, 10, 15, and 30-minute horizons.
- **The “Complexity Tax”:** The full DCRNN model using sensor adjacency performed significantly worse than the LSTM, with Validation Loss nearly **2x-3x higher** in some trials.

**Unexpected Result:** The “Speed-only” DCRNN (no graph adjacency) outperformed the full Graph DCRNN, suggesting the spatial data was hindering, not helping, our specific implementation.

# Discussion

**Adjacency Matrix Limitations:** The pre-defined sensor graph may not accurately reflect true traffic dynamics (e.g., functional connectivity vs. physical distance).

## Hyperparameter Sensitivity:

- Graph operations increase model complexity, making convergence much harder.
- Parameters tuned for sequence models (LSTM) do not transfer directly to diffusion convolutional architectures.

## Architecture Isolation:

- Comparing **LSTM vs. DCRNN (No-Adj)** isolates the architecture.
- Since DCRNN (No-Adj) was competitive but still slightly worse, the deficit lies in the *graph construction*, not necessarily the diffusion mechanism.

# Conclusion

**Simplicity remains competitive:** For short-term forecasting (5-15 mins), a strong temporal baseline (LSTM) is efficient and robust.

**Complexity is not a guarantee:** Simply switching to a graph-based architecture does not guarantee performance gains without rigorous tuning and graph definition.

- Our DCRNN with adjacency did not reach performance reported in prior work
- Results underscore the importance of careful graph design and replication of reference setups

## Future Work:

- **Dynamic Graph Learning:** Instead of a static distance matrix, let the model learn the graph structure (e.g., Graph WaveNet).
- **Hybrid Tuning:** Investigate separate hyperparameter optimization strategies specifically for the diffusion layers.

# Acknowledgments

Special Thanks to:

Our Team:

- Elyann Soto
- Krishna Sharma
- Martin Gasevski
- Sterling Walker

Our Data:

- California Department of Transportation - Performance Measurement System (PeMS)

DCRNN Reference Implementation:

- <https://github.com/liyaguang/DCRNN>

*Thank  
you!*